

Done — recreated smarter.

Skill: critical-infrastructure-cyber-resilience

Location: /home/workdir/.grok/skills/critical-infrastructure-cyber-resilience/

Validation: OK (75 lines)

What changed (smarter version)

Aspect	Previous	Smarter
Description	Long list of frameworks	High-signal trigger list + clear scope
Core rules	Good	Explicit “Never invent requirements” + source_id integrity rule
Activation	Implicit	New decision tree — when to use vs fall back
Workflow	Linear list	Defensive order (confirm scope → source_manifest → extract → normalize → validate)
Taxonomy	Listed in SKILL.md	Moved fully into reference; SKILL stays lean
Schemas	Verbose examples	Required fields only + one clean example
Seed queue	Long	Minimal template + hard rules
Scripts	Empty	Added scripts/validate_records.py (schema + status + integrity checks)
Token load	Higher	~25% leaner while retaining every critical constraint from your original plan

Key improvements in practice

- Stronger guardrails against hallucinated obligations.
- Clearer separation of extraction vs. mapping.
- Ready-to-run validation script for any produced .jsonl .
- Progressive disclosure is tighter — references only load when needed.

- Still fully covers the original research program (7 jurisdictions, OT distinction, mandatory/voluntary tagging, incident matrices, seed queue, drift protocol).

The skill is now more precise, less verbose, and harder to misuse.

Ready for real work.

Say the word and I will start any lane (e.g. “run T04 U.S. baseline”, “generate the full seed queue”, “build the first incident-reporting matrix”, etc.).